

One convolutional-network training step. Losses match when they differ by at most 0.02.

program	runs to completion?	loss	gradient norm	parameter update	completion + loss verdict
Source program the reference	yes	1.7805	0.2834	reference	--
Correct port no fault	yes	1.7805	0.2834	matches	accept
Missing operation no implementation	no	--	--	--	reject
Wrong reduction averaged, not summed	yes	2.5009	0.6879	differs	reject
No backward pass no gradients	yes	1.7805	--	no update	accept

That verdict accepts both shaded rows, though their gradient and update evidence differs.

One of the two never updates its parameters, and that check cannot tell them apart.